

Data Communication and Computer Networks

3. Application Layer PART-B

Dr. Aiman Hanna

Department of Computer Science & Software Engineering
Concordia University, Montreal, Canada

These slides has mainly been extracted, modified and updated from original slides of :
Computer Networking: A Top Down Approach, 6th edition Jim Kurose, Keith Ross
Addison-Wesley, 2013

Additional materials have been extracted, modified and updated from:
Understanding Communications and Networking, 3e by William A. Shay 2005

Copyright © 1996-2013 J.F Kurose and K.W. Ross
Copyright © 2005 William A. Shay
Copyright © 2019 Aiman Hanna
All rights reserved

DNS: domain name system

people: many identifiers:

- SSN, name, passport #

Internet hosts, routers:

- IP address (32 bit) - used for addressing datagrams
- “name”, e.g., `www.yahoo.com` - used by humans

Q: how to map between IP address and name, and vice versa ?

Domain Name System:

1) *distributed database*

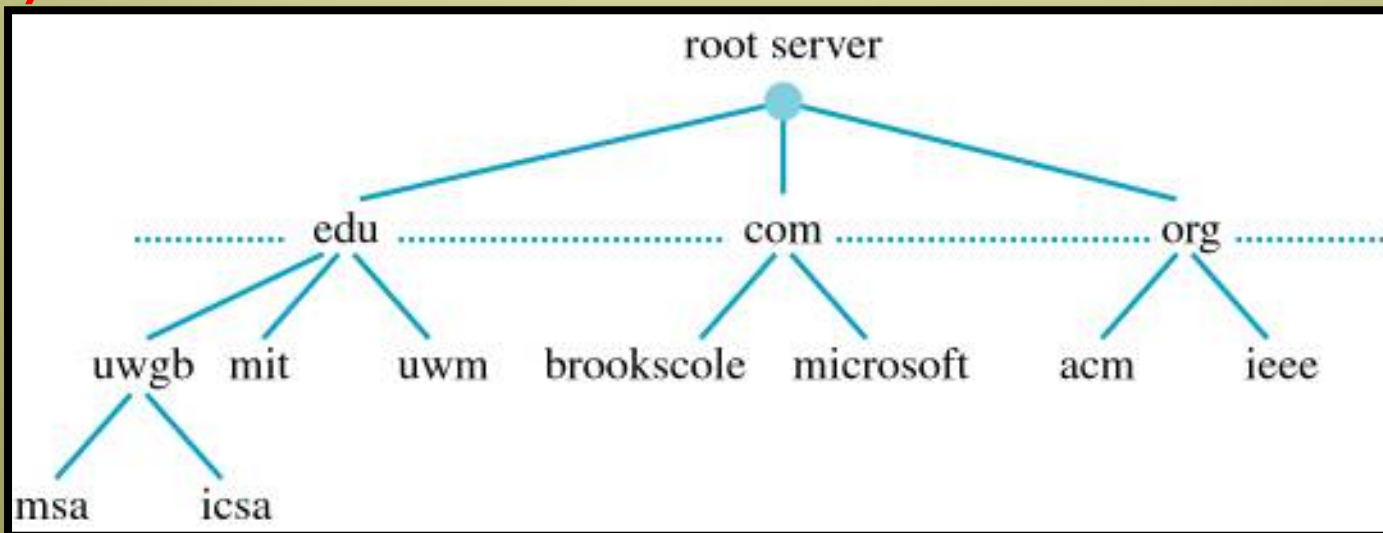
implemented in hierarchy of many *DNS servers*

2) *application-layer protocol*: which allows hosts to query the database

- name servers communicate to *resolve* names (address/name translation)
- note: core Internet function, implemented as application-layer protocol
- complexity at network's “edge”; possibly a substantial delay

DNS: domain name system

- ❖ To facilitate the lookup process, servers are organized into zones
- ❖ A request for a textual name escalates, and may also go down, until the IP address is found (or the process fails)
- ❖ *why not centralize DNS?*



DNS: services, structure

- ❖ **host aliasing**: specify domain name as an alias of another
 - i.e. point **ftp.example.com** and **www.example.com** (an FTP server and a webserver running on two ports from a single IP address) to **example.com**
 - such domain names are referred to a Canonical Names (**CNAMEs**); the pointed-to domain name is an **alias**
 - a CNAME must point to another domain name; never to an IP address. In turn that pointed-to domain name points to the IP address
 - if the IP of the server ever changes, only one update to the DNS record is required for all CNAMEs

DNS: services, structure

❖ mail server aliasing


- e-mail addresses need to be resolved as well; however the name e-mail server may be CNAME
- i.e. an e-mail directed to linda@example.com may seem to require resolving example.com; however the mail server at example.com may actually be a CNAME (i.e relay.west-coast.example.com)
- DNS can be invoked by e-mail application to obtain the CNAME of an alias hostname, as well as the IP address of the host

DNS: services, structure

❖ **load distribution**

- heavy-loaded sites are often replicated over multiple servers
- consequently, a set of IP addresses is associated with one hostname
- DNS database contains the set of these IP addresses
- the entire set is returned upon a query from a DNS client
- client often sends its HTTP request to the first IP listed in the returned list
- DNS rotations attempts to distribute the load over the replicated servers by rotating the order of the returned IP addresses

DNS: a distributed, hierarchical database

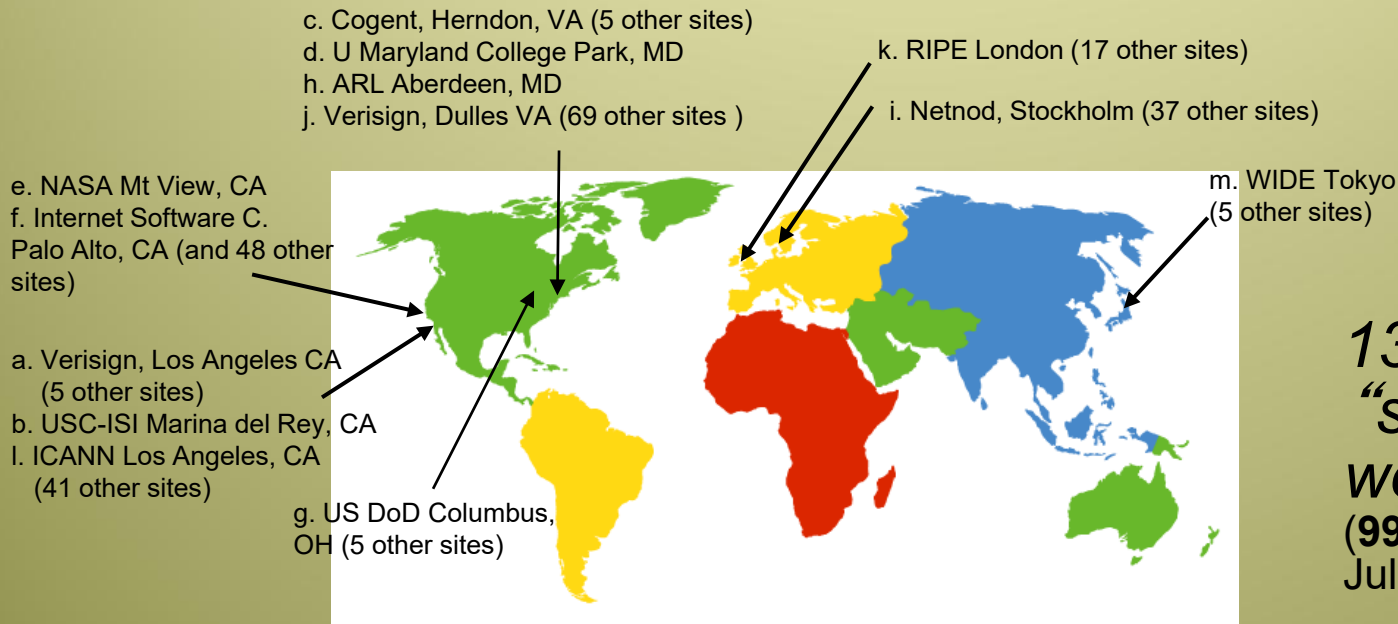


client wants IP for www.amazon.com; 1st approx:

- ❖ client queries root server to find com DNS server
- ❖ client queries .com DNS server to get amazon.com DNS server
- ❖ client queries amazon.com DNS server to get IP address for www.amazon.com

DNS: root name servers

- ❖ contacted by local name server that can not resolve name
- ❖ root name server:
 - contacts authoritative name server if name mapping not known
 - gets mapping
 - returns mapping to local name server



*13 root name
“servers”
worldwide
(997 instances as of,
July 11, 2019)*

TLD & authoritative servers

top-level domain (TLD) servers:

- responsible for com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g.: uk, fr, ca, jp
- Verisign/Network Solutions maintains servers for .com TLD
- Educause for .edu TLD

authoritative DNS servers:

- organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
- can be maintained by organization or service provider

Local DNS name server

- ❖ does not strictly belong to DNS hierarchy
- ❖ each ISP (residential ISP, company, university) has one
 - also called “default name server”
- ❖ when host makes DNS query, query is sent to its local DNS server
 - has local cache of recent name-to-address translation pairs (but may be out of date!)
 - acts as proxy, forwards query into hierarchy

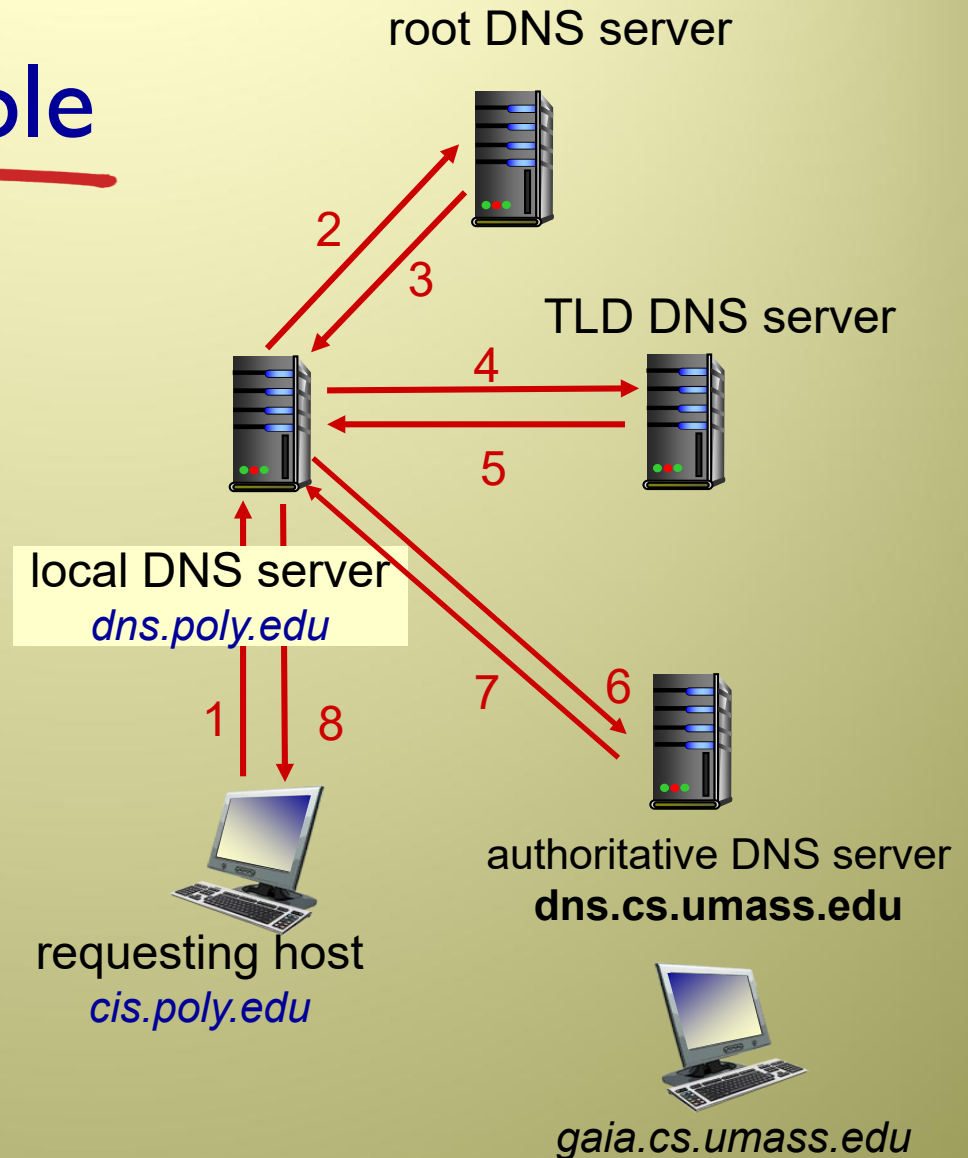
DNS name resolution example

- ❖ host at cis.poly.edu wants IP address for gaia.cs.umass.edu

Iterative query:

- ❖ contacted server replies with name of server to contact
- ❖ “I don’t know this name, but ask this server”

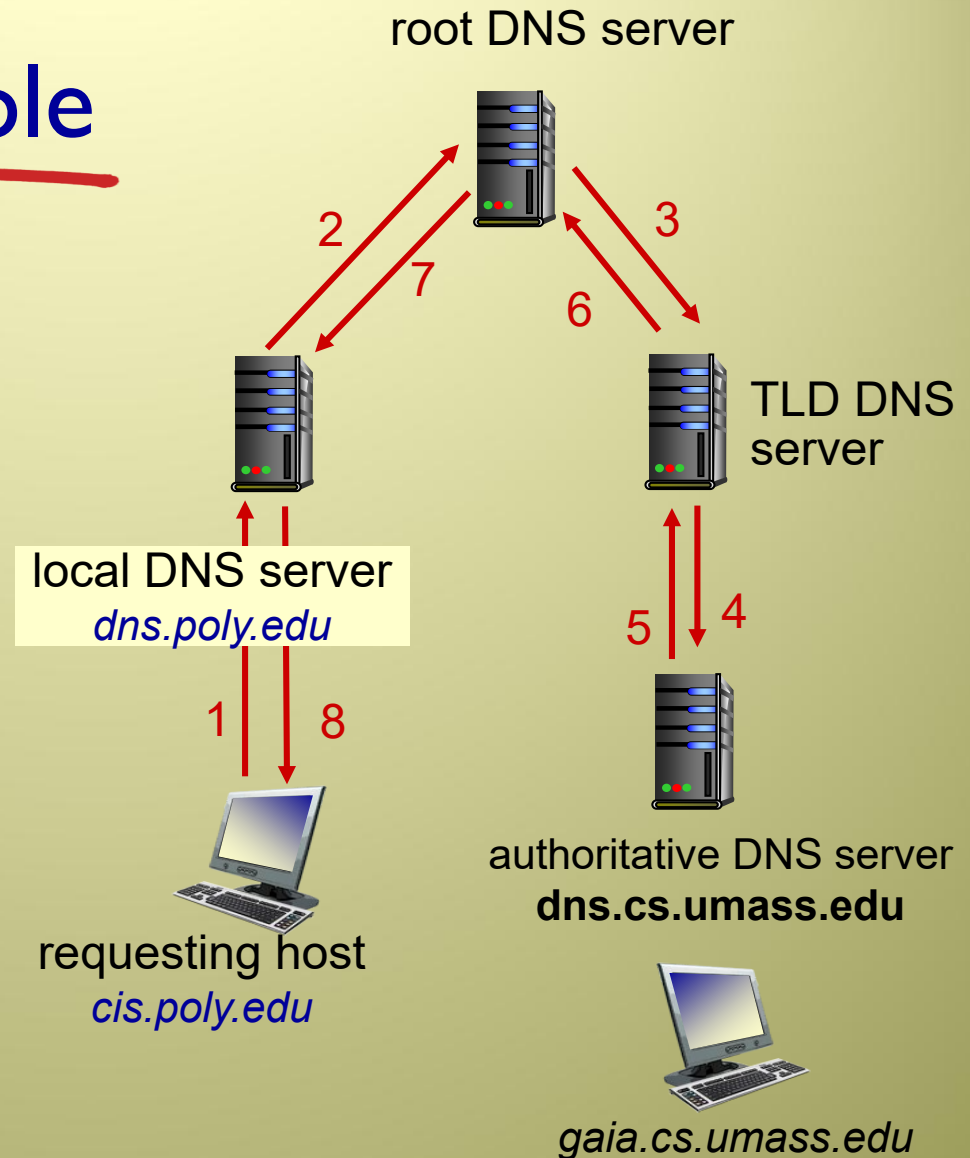
Note: call from cis.poly.edu is recursive; rest are iterative



DNS name resolution example

recursive query:

- ❖ puts burden of name resolution on contacted name server
- ❖ heavy load at upper levels of hierarchy?



DNS: caching, updating records

- ❖ a critically important feature of DNS
- ❖ improve delay performance & reduce DNS messages in the Internet
- ❖ once (any) name server learns of an address mapping, it *caches* it
 - cache entries timeout (disappear) after some time (TTL)
 - TLD servers typically cached in local name servers
 - thus root name servers not often visited

DNS: caching, updating records

- ❖ cached entries may be *out-of-date* (best effort name-to-address translation!)
 - if name host changes IP address, may not be known Internet-wide until all TTLs expire
- ❖ update/notify mechanisms proposed IETF standard
 - RFC 2136

DNS records

DNS: distributed db storing resource records (RRs)

RR format: (Name, Value, Type, TTL)

type=A

- name is hostname (e.g, relay.west-coast.example.com)
- **value** is IP address

Example: (relay1.bar.foo.com, 145.46.93.9, A)

type=NS

- **name** is domain (e.g., foo.com)
- **value** is hostname of authoritative name server for this domain

Example: (foo.com, dns.foo.com, NS)

type=CNAME

- name is alias name for some “canonical” (the real) name
- www.ibm.com is really servereast.backup2.ibm.com
- value is the canonical name

Example: (foo.com, relay1.bar.foo.com, CNAME)

type=MX

- name is an alias host name
- value is a canonical name of mailserver associated with name

Example: (foo.com, mail.bar.foo.com, MX)

DNS records

- each DNS reply may carry one or **more** RRs

Example

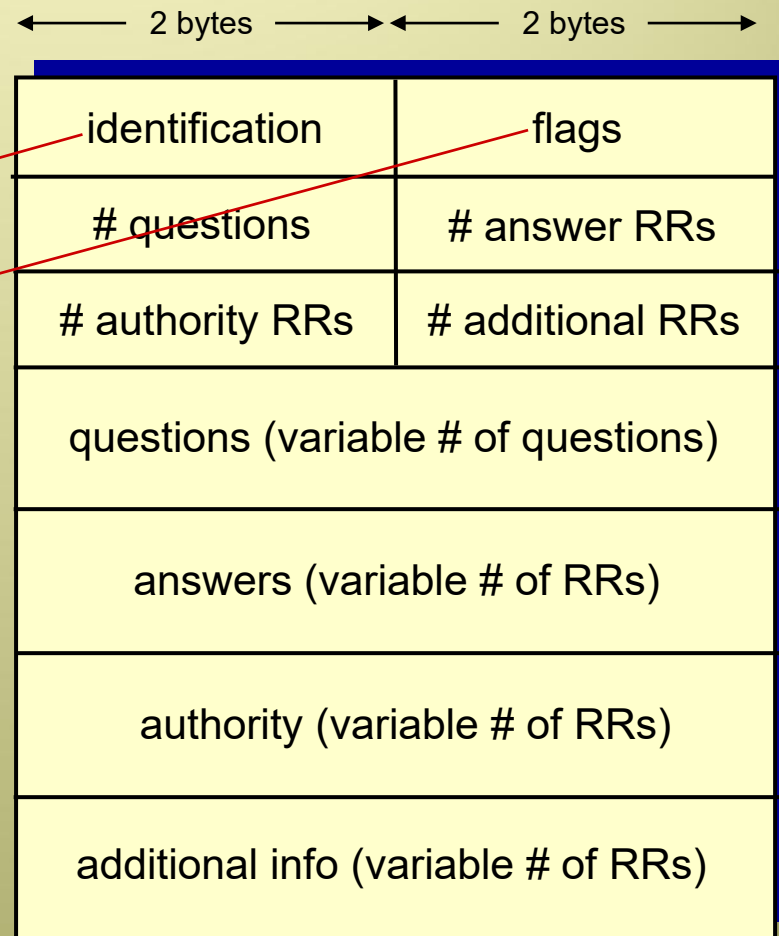
- ❖ Assume the DNS server is an authoritative server for a particular hostname
 - server will return Type A RR
e.g., (gaia.cs.umass.edu, 145.119.96.44, A)
 - Note: This type may be returned by any non-authoritative server as well (i.e. from a previously cached record)
 - ❖ Assume the DNS server is NOT an authoritative server for a particular hostname
 - server may return Type NS RR
e.g., (umass.edu, dns.umass.edu , NS)
- As well as:
- e.g., (dns.umass.edu, 145.119.96.28, A)

DNS protocol, messages

- ❖ *query* and *reply* messages, both have same *message format*

msg header

- ❖ **identification:** 16 bit #
for query, reply to query
uses same #
- ❖ **flags:**
 - 1-bit query or reply
 - 1-bit recursion
desired
 - 1-bit recursion
available
 - 1-bit reply is
authoritative



DNS protocol, messages

name, type fields for a query

i.e. host address associated with a name (type A), or the mail server for a name (type MX)

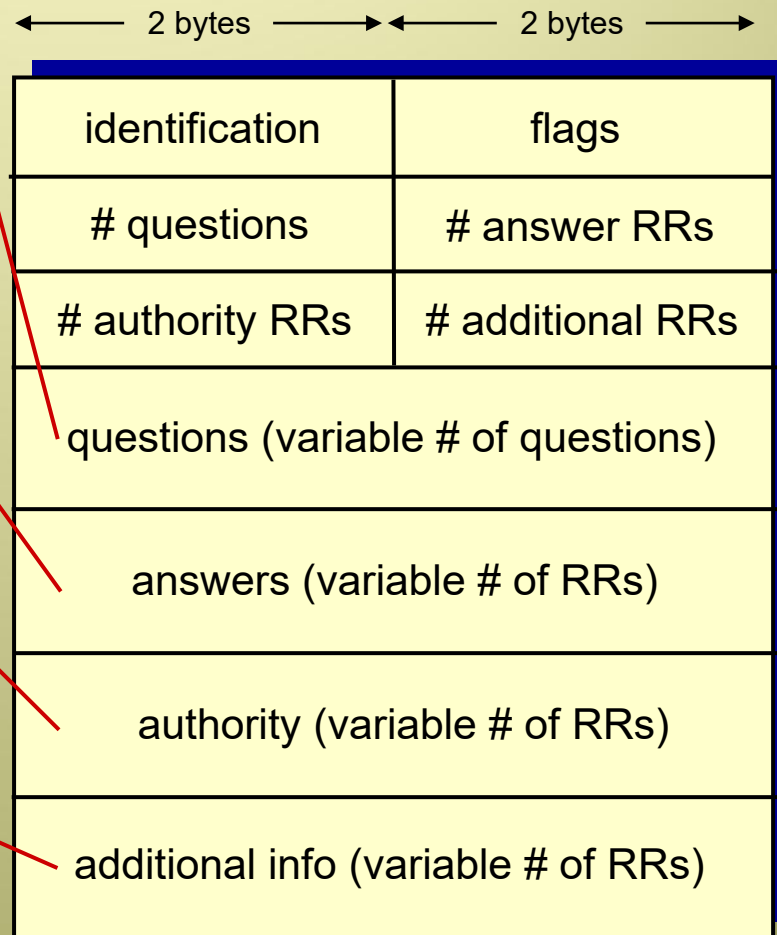
RRs in response to query

P/S: A reply may contain multiple RRs (i.e. for replicated web servers)

records for authoritative servers

additional “helpful” info that may be used

i.e. the answers section of a reply to an MX query include a RR providing the CNAME of the mail server. In the additional section may contain a type A RR providing the IP address of this CNAME



Inserting records into DNS

- ❖ example: new startup “Network Utopia”
- ❖ register name networkutopia.com at *DNS registrar* (e.g., Network Solutions)
 - provide names, IP addresses of authoritative name server (primary and secondary)
 - *registrar* inserts two RRs into .com TLD server:
(networkutopia.com, dns1.networkutopia.com, NS)
(dns1.networkutopia.com, 212.212.212.1, A)
- ❖ enter needed records into your authoritative DNS server:
 - type A record for www.networkutopia.com;
 - type MX record for mail.networkutopia.com

Attacking DNS

DDoS attacks

- ❖ Bombard root servers with traffic
 - i.e. large-scale attack on October 21, 2002 using massive *ping* messages to the 13 root servers
 - Not successful to date
 - Traffic Filtering (configured to block *ping* messages)
 - Local DNS servers cache IPs of TLD servers, allowing root server bypass
- ❖ Bombard TLD servers
 - Potentially more dangerous
 - i.e. send a massive number of DNS queries; which will not be filtered
 - Severity can still be mitigated (at least partially) by caching in local DNS servers

Attacking DNS

Redirect attacks

- ❖ Man-in-middle
 - Intercept queries, then send bogus replies
- ❖ DNS poisoning
 - Send bogus replies to DNS server, which caches these incorrect replies
 - Connections to the intended site can then be directed to the attacker's site

Exploit DNS for DDoS (not directly an attack on DNS but on a targeted host)


- ❖ Send queries to many authoritative DNS servers with spoofed source address as target IP
- ❖ Requires amplification (responses are much larger than queries), which overwhelm the target
 - Notice that this is done without the need for the attacker to generate much of its own traffic

Pure P2P architecture

- ❖ no always-on server
- ❖ arbitrary end systems directly communicate
- ❖ peers are intermittently connected and change IP addresses

examples:

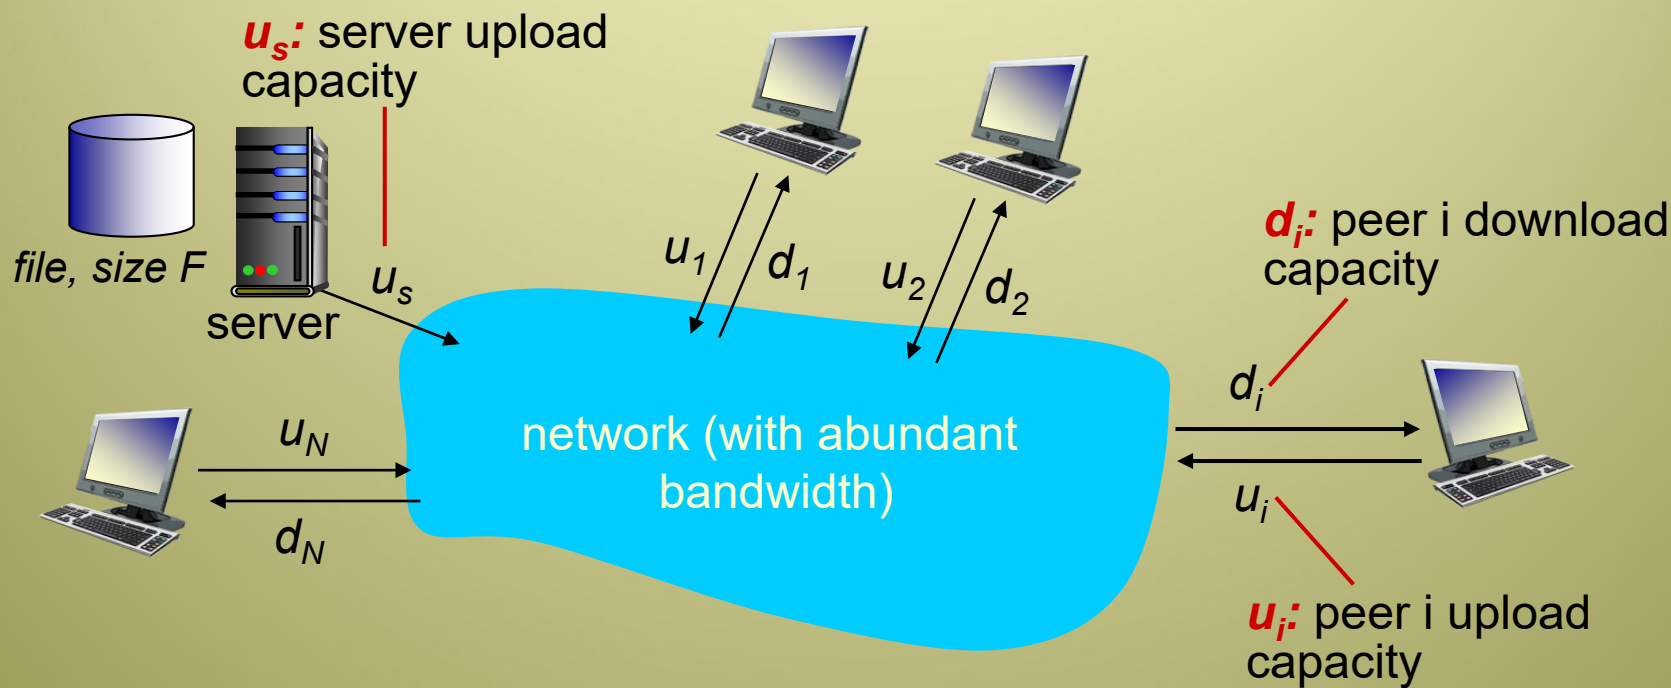
- file distribution (BitTorrent)
- Streaming (KanKan)



File distribution: client-server vs P2P


Question: how much time to distribute file (size F) from one server to N peers (clients)?

- peer upload/download capacity is limited resource



File distribution time: client-server

- ❖ **server transmission:** must sequentially send (upload) N file copies:
 - total transmission: FN bits
 - time to send one copy: F/u_s
 - time to send N copies is: NF/u_s
- ➔ distribution time is, at least:
 NF/u_s




- ❖ **client:** each client must download file copy
- ❖ Let d_{min} denote the download rate of the peer with the slowest download rate
 - $d_{min} = \min\{d_1, d_2, \dots, d_N\}$
 - min-client download time: F/d_{min}
- ➔ distribution time is, at least: F/d_{min}

File distribution time: client-server

→ To send, distribution time is, at least: NF/u_s

→ To receive, distribution time is, at least: F/d_{min}

❖ Let D_{cs} denote distribution for client-server architecture



$$D_{cs} \geq \max\{ NF/u_s, F/d_{min} \}$$


increases linearly in N

$$D_{cs} = \max\{ NF/u_s, F/d_{min} \}$$

Lower-bound (best case) distribution time

File distribution time: P2P

- ❖ each peer can assist the server in distributing the file
- ❖ **server transmission:** must upload at least one copy
 - time to send one copy: F/u_s
- ❖ **client:** each client must download file copy
 - min client download time: F/d_{\min}
 - distribution time is, at least: F/d_{\min}
- ❖ Total upload capacity of the **system as a whole:**
$$u_{\text{total}} = u_s + u_1 + u_2 + u_3 + \dots + u_N$$
- ❖ System still needs to deliver F bits to N peers, totaling to **NF** bits, which cannot be done at a rate faster than **u_{total}**




→ distribution time is, at least: **NF/u_{total}**

File distribution time: P2P

→ server time to send one copy: F/u_s

→ client distribution time is, at least:
 F/d_{min}

→ System-wide distribution time is, at least: NF/u_{total}




❖ Let D_{P2P} denote distribution for P2P architecture

$$D_{P2P} \geq \max\{ F/u_s, F/d_{min}, NF/(u_s + \sum_{i=1}^N u_i) \}$$

increases linearly in N ...

... but so does this, as each peer brings service capacity

File distribution time: P2P




❖ Lower-bound (actual minimum distribution time) for P2P distribution can then be given as:

$$D_{P2P} = \max\{ F/u_s, F/d_{min}, NF/(u_s + \sum_{i=1}^N u_i) \}$$

Lower-bound (best case) distribution time

Client-server vs. P2P: example

client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{min} \geq u_s$




P2P file distribution: BitTorrent

- ❖ file divided into 256KByte chunks
- ❖ peers in torrent send/receive file chunks


tracker: tracks peers participating in torrent

torrent: group of peers participating in the distribution (exchanging chunks of a file)



P2P file distribution: BitTorrent

- ❖ peer joining torrent:
 - has no chunks, but will accumulate them over time from other peers
 - registers with **tracker** to get list of peers, connects to subset of peers (“neighbors”)
- ❖ while downloading, peer uploads chunks to other peers
- ❖ peer may change peers with whom it exchanges chunks
- ❖ **churn**: peers may come and go
- ❖ once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent



BitTorrent: requesting, sending file chunks

requesting chunks:


- ❖ at any given time, different peers have different subsets of file chunks
- ❖ periodically, Alice asks each peer for list of chunks that they have
- ❖ The provided list from each of the peers allows Alice to make 2 important decisions:
 - which chunks should she request first (**rarest first**)
 - to which of her neighbors should she send requested chunks

sending chunks: tit-for-tat

- ❖ Alice sends chunks to the four peers currently sending her chunks *at highest rate*
 - other peers are **choked** by Alice (do not receive chunks from her)
 - re-evaluate top 4 every 10 secs
- ❖ every 30 secs: randomly select another peer, starts sending chunks
 - “**optimistically unchoke**” this peer
 - newly chosen peer may join top 4

BitTorrent: tit-for-tat

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



*higher upload rate: find
better trading partners, get
file faster !*