

Data Communication and Computer Networks

7. Network Layer PART-B

Dr. Aiman Hanna

Department of Computer Science & Software Engineering
Concordia University, Montreal, Canada


These slides have mainly been extracted, modified and updated from original slides of :
Computer Networking: A Top Down Approach, 6th edition Jim Kurose, Keith Ross
Addison-Wesley, 2013

Additional materials have been extracted, modified and updated from:
Understanding Communications and Networking, 3e by William A. Shay 2005

Copyright © 1996-2013 J.F Kurose and K.W. Ross
Copyright © 2005 William A. Shay
Copyright © 2019 Aiman Hanna
All rights reserved

Wide Area Networks (WANs)


- ❖ **Wide Area Networks, WANs**, are large networks that may span the globe
- ❖ WANs have a generalized topologies and require more sophisticated techniques to perform what they are expected to



Generalized Network Topology

Routing Tables

- ❖ these table do not normally specify the entire route
- ❖ rather, they specify the next node in a route to a specified destination, and the cost to get there



Network Example

Routing Tables


- ❖ Example: partial tables for the routers shown in the example network

DESTINATION	NEXT NODE	COST	DESTINATION	NEXT NODE	COST	DESTINATION	NEXT NODE	COST
B	B	2	D	D	5	F	F	2
C	C	1	E	E	3			
D	C	5	F	E	5			
E	B	5						
F	B	7						

(a) Partial routing table for node A (b) Partial routing table for node B (c) Partial routing table for node E

Partial Routing Tables for A, B & E

Graph abstraction




graph: $G = (N, E)$ - Nodes and Edges

N = set of nodes/routers = $\{ u, v, w, x, y, z \}$

E = set of links/edges

= $\{ (u,v), (u,x), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z) \}$

Graph abstraction: costs



$c(x, x') = \text{cost of link } (x, x')$

e.g., $c(w, z) = 5$

cost could always be 1, or
inversely related to bandwidth,
or inversely related to
congestion

cost of path $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

key question: what is the least-cost path between u and z ?

routing algorithm: algorithm that finds that least cost path

Routing algorithm classification

Q: global or decentralized information?

global:

- ❖ all routers have complete topology, link cost info
- ❖ “link state” algorithms

decentralized:

- ❖ router knows physically-connected neighbors, link costs to neighbors
- ❖ iterative process of computation, exchange of info with neighbors
- ❖ “distance vector” algorithms

Q: static or dynamic?

static:

- ❖ routes change slowly over time

dynamic:

- ❖ routes change more quickly
 - periodic update
 - in response to link cost changes

Routing algorithm classification

Centralized Routing

- ❖ all interconnection information is generated and maintained at a single central station
- ❖ that station then broadcasts this information to all network nodes, so they can define their own routing tables
- ❖ one way to maintain the central information is through **routing matrix**

Routing algorithm classification

Centralized Routing

		A	B	C	D	E	F
source node	A	—	B	C	C	B	B
	B	A	—	A	D	E	E
	C	A	A	—	D	E	F
	D	C	B	C	—	F	F
	E	B	B	C	F	—	F
	F	E	E	C	D	E	—

Routing Matrix for the Example Network

Routing algorithm classification

Distributed Routing

- ❖ no central control; each node must determine and maintain routing information independently
- ❖ this can be achieved by knowing the neighbors, cost to get to a neighbor and the cost from that neighbor to a destination
- ❖ more complex than centralized since each node must communicate with each of its neighbors instead of just one central station
- ❖ the complexity is driven also by the fact that the devices have a very limited knowledge of the entire network

Routing algorithm classification

Static Routing

- ❖ once the node determines the routing table, it does not change it
- ❖ in other words, the initial cheapest path may not really be the cheapest path after sometime, yet it is always considered as the cheapest path
- ❖ the assumption here is that the conditions that led to the initial definition of the tables do not change
- ❖ sometimes, this is a reasonable assumption, when?

Routing algorithm classification

Adaptive Routing

- ❖ allows the network to respond to changes and update its routing tables accordingly
- ❖ if the cost of the route to be used changed, then adapt and use another route
- ❖ could this lead to serious problems sometime?
- ❖ in general, it is difficult to implement adaptive routing efficiently, why?

A Link-State Routing Algorithm

Dijkstra's algorithm

- ❖ sometimes called the ***Shortest-path Algorithm*** or ***Forward Search Algorithm***
- ❖ centralized, static algorithm, however it can be made adaptive by executing it periodically
- ❖ a node executing this algorithm is required to know the link costs among the nodes
- ❖ each node executes this algorithm to determine the cheapest route to each network node

A Link-State Routing Algorithm

Dijkstra's algorithm

- ❖ net topology, link costs known to all nodes
 - accomplished via “link state broadcast”
 - all nodes have same info
- ❖ computes least cost paths from one node (‘source’) to all other nodes
 - gives *forwarding table* for that node
- ❖ iterative: after k iterations, know least cost path to k destinations

notation:

- ❖ $c(x,y)$: link cost from node x to y; $= \infty$ if not direct neighbors
- ❖ $D(v)$: current value of cost of path from source to dest. v
- ❖ $p(v)$: prior (predecessor) node along path from source to v
- ❖ S : set of nodes whose least cost path definitively known

Dijkstra's Algorithm

1 **Initialization:**

2 $S = \{u\}$

3 for all nodes v

4 if v adjacent to u

5 then $D(v) = c(u,v)$

6 else $D(v) = \infty$

7

8 **Loop**

9 find w not in S such that $D(w)$ is a minimum

10 add w to S

11 update $D(v)$ for all v adjacent to w and not in S :

12 **$D(v) = \min(D(v), D(w) + c(w,v))$**

13 /* new cost to v is either old cost to v or known

14 shortest path cost to w plus cost from w to v */

15 **until all nodes in S**


Dijkstra's Algorithm

- ❖ Iterative procedure.

C: Set of nodes to which cheap route is known,

S: Set of nodes directly connected to a node in C,

X: A node in S to where
cheap path found



Distances yet to be determined

Distances to potential vertices

Potential vertex with shortest distance

X	Selected vertex to move to C
-----	--------------------------------


Different routes to u .

Is there a cheaper route through X to other vertices in S ? If so, update distances to these vertices.

Dijkstra's Algorithm

- ❖ In details, the algorithm performs the repeated moving operations from S to C as follows:
 - Move one vertex from S to C . Let us refer to that vertex as X . This vertex, X , must satisfy the following:
 - 1) It is either directly connected to v or directly connected to another vertex that is already in C
 - 2) Has the smallest distance to v among all the potential vertices that are considered for movement to C
 - Whenever the vertex X is moved to C , its distance is the smallest distance (shortest path) to v
 - Additionally, once this vertex X is chosen, move it to C and recheck all distances of all vertices that have direct connection to X and are NOT yet in C . If a smaller value than what we already have is found, update this value
 - In each step, we also update how to go to these vertices (that is, through which vertex); we refer to this vertex as the **prior function**
 - The steps are repeated until all the graph vertices are moved to C . The final obtained distances represent the shortest paths, and the prior functions indicate the direction from v (through which vertex) to obtain these shortest paths to each of the other vertices in the graph

Dijkstra's Algorithm - Example



Network and Associated Connection Costs

Dijkstra's Algorithm - Example

Cost function Prior function


	C	Potential elements of S	X	A	B	C	D	E	F	A	B	C	D	E	F
0	{}	{A}	A	0	∞	∞	∞	∞	∞	A	-	-	-	-	-
1	{A}	{B,C}	C	0	2	1	∞	∞	∞	A	A	A	-	-	-
2	{A, C}	{B,D,E,F}	B	0	2	1	4	7	8	A	A	A	C	C	C
3	{A,B,C}	{D,E,F}	D	0	2	1	4	6	8	A	A	A	C	B	C
4	{A,B,C,D}	{E,F}	E	0	2	1	4	6	6	A	A	A	C	B	D
5	{A,B,C,D,E}	{F}	F	0	2	1	4	6	6	A	A	A	C	B	D

19

Note: See comments attached to this slide for detailed information on how the operations progressed.

Dijkstra's algorithm: example

resulting shortest-path tree from A:



resulting forwarding table in A:

destination	link
B	(A,B)
C	(A,C)
D	(A,C)
E	(A,B)
F	(A,C)

Hierarchical routing

our routing study thus far - idealization

- ❖ all routers identical
- ❖ all run the same routing algorithm
- ❖ network “flat”

... *not* true in practice

scale: with 600 million destinations:

- ❖ cannot store all dest' s in routing tables!
- ❖ routing table exchange would swamp links!

administrative autonomy

- ❖ internet = network of networks
- ❖ each network admin may want to control routing in its own network, while still being able to connect to other networks

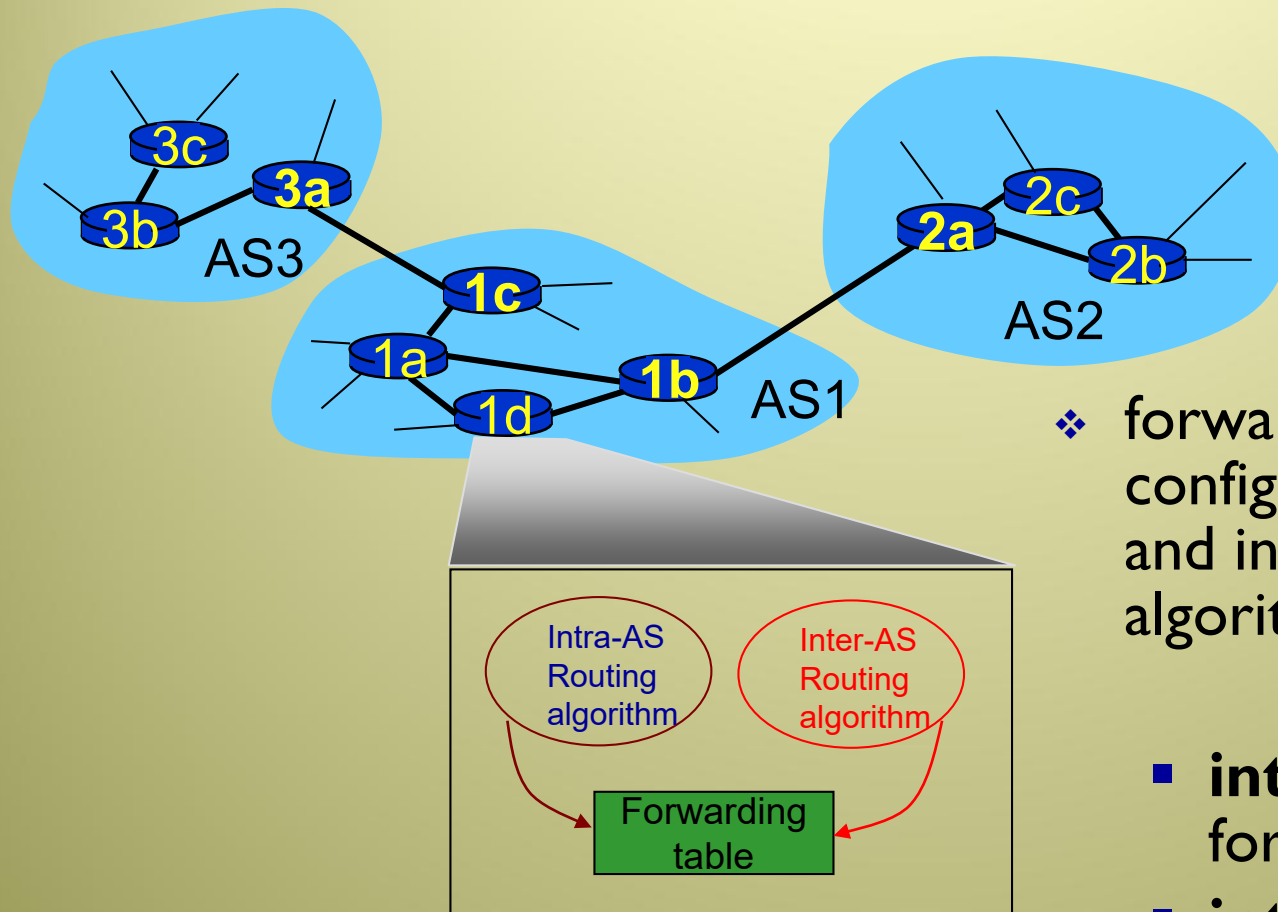
Hierarchical routing

- ❖ organize routers into regions, “**autonomous systems**” (AS)
- ❖ routers in same AS:
 - typically under the same admin control (i.e. operated by same ISP, or belonging to same company)
 - run same routing protocol
 - “**intra-AS**” **routing** protocol
 - routers in different AS can run different **intra-AS routing protocol**

gateway router:

- ❖ an “edge” of its own AS
- ❖ has the task of forwarding packets to destination outside its AS
- ❖ has link to router in another AS

Interconnected ASes



❖ forwarding table configured by both intra- and inter-AS routing algorithm

- **intra-AS** sets entries for internal dests
- **inter-AS & intra-AS** sets entries for external dests


Inter-AS tasks

- ❖ suppose router in AS1 receives datagram destined outside of AS1:
 - router should forward packet to gateway router, but which one?

AS1 must:


1. learn which destds are reachable through AS2, which through AS3
2. propagate this reachability info to all routers in AS1

job of inter-AS routing!




Example: setting forwarding table in router 1d

- ❖ suppose AS1 learns (via inter-AS protocol) that subnet **x** reachable via AS3 (gateway 1c), but not via AS2
 - inter-AS protocol propagates reachability info to all internal routers
- ❖ router 1d determines from intra-AS routing info that its interface **I** is on the least cost path to 1c
 - installs forwarding table entry **(x,I)**



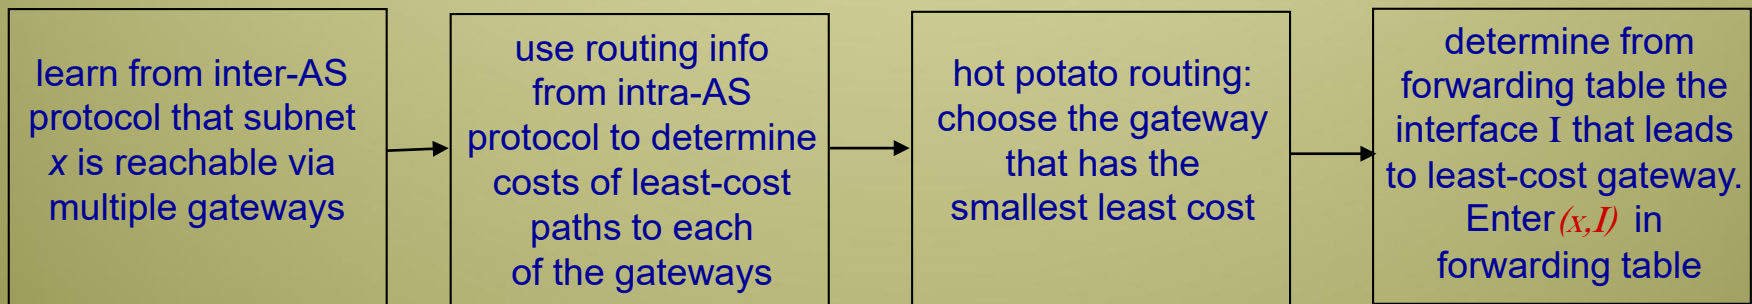
Example: choosing among multiple ASes

- ❖ now suppose AS1 learns from inter-AS protocol that subnet **x** is reachable from AS3 *and* from AS2.
- ❖ to configure forwarding table, router 1d must determine which gateway it should forward packets towards for dest **x**
 - this is also job of inter-AS routing protocol!




Example: choosing among multiple ASes

- ❖ if more than one gateway can be used, then the router must determine which one packets must be forwarded
- ❖ one common approach is the: **hot-potato routing**
- ❖ **hot potato routing:**
 - AS gets rid of the packet as quickly as possible
 - send packet towards closest of the two (or the many) routers



Broadcast routing

- ❖ deliver packets from source to **all** other nodes
- ❖ source duplication (**N-way-unicast**) is inefficient:




- ❖ source duplication: how does source determine recipient addresses?

In-network duplication

- ❖ *uncontrolled flooding*: when node receives broadcast packet, sends copy to all neighbors
 - problems: cycles & broadcast storm
- ❖ *controlled flooding*: node only broadcasts pkt if it hasn't broadcast same packet before
 - node keeps track of packet ids already broadcasted
 - or reverse path forwarding (RPF): only forward packet if it arrived on shortest path between node and source
- ❖ *spanning tree*:
 - no redundant packets received by any node

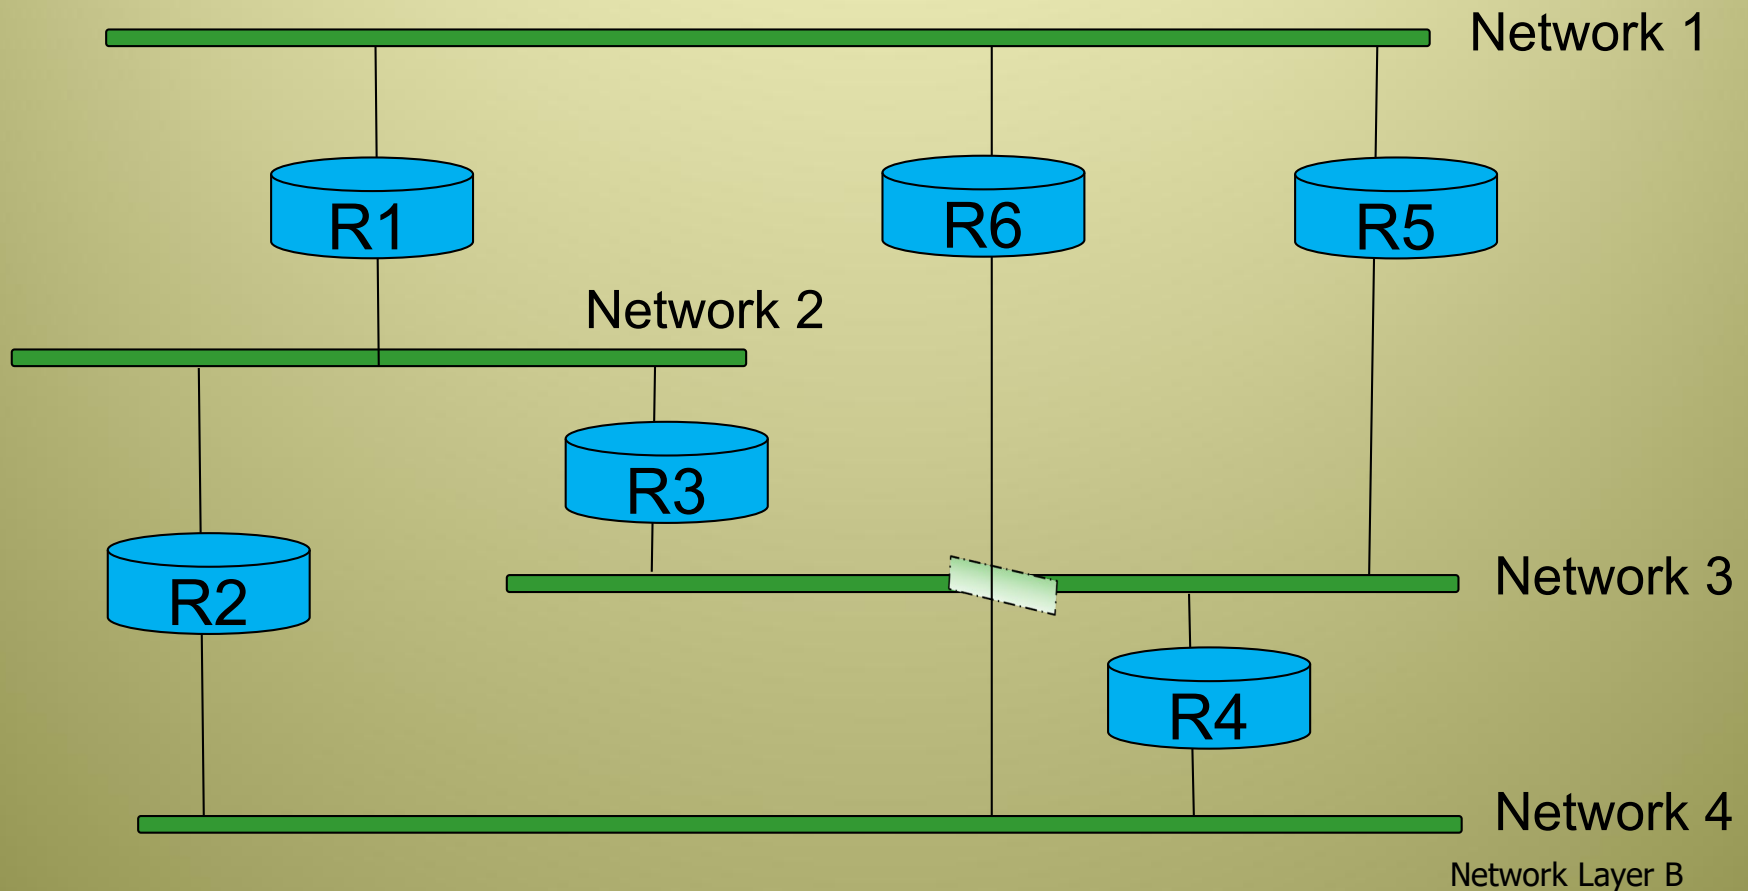
Uncontrolled flooding



- ❖ What happens when a device on network 1 sends a packet to another on network 2?

Uncontrolled flooding

- ❖ no multiple routers in between any two networks. Is there still a problem now?



Sequence-number controlled flooding


- ❖ source node puts its address (or a unique identifier), as well as a **broadcast sequence number** into the broadcast packet, then
- ❖ sends a copy to each of its neighbors
 - problems: potential cycles & broadcast storm
- ❖ each node maintains a list of the source address/sequence number for each broadcast packet that it received, duplicated, and forwarded
- ❖ if a received packet is in that list, the packet is dropped; if not, the packet is duplicated and forwarded to all neighbors, except the one it came from

Reverse Path Forwarding (RPF) controlled flooding

- ❖ when a router receives a broadcast packet from a source address, it forwards it to all its neighbors (except where it came from) ONLY IF:
 - Packet arrived on the link that is on its own shortest unicast path back to the source
 - Otherwise, packet is discarded
- ❖ packets from any unexpected paths are dropped since the router knows that it either will receive, or has already received, a copy of this packet on the link that is on its own shortest path back to the sender

Reverse Path Forwarding (RPF) controlled flooding

- ❖ Example: packets are sent from source A.
 - least-cost paths are represented by thick lines




→ ⊗ Packet will not be forwarded beyond receiving router

→ Packet will be forwarded


— Thick lines represent least cost paths from receiver to source A

Spanning tree

- ❖ both sequence-number and RPF controlled flooding do not completely avoid the transmission of redundant broadcast packets
- ❖ first construct a spanning tree
- ❖ nodes then forward/make copies only along spanning tree




(a) broadcast initiated at A




(b) broadcast initiated at D

Spanning tree: creation

- ❖ center node
- ❖ each node sends unicast join message to center node
 - message forwarded until it arrives at a node already belonging to spanning tree



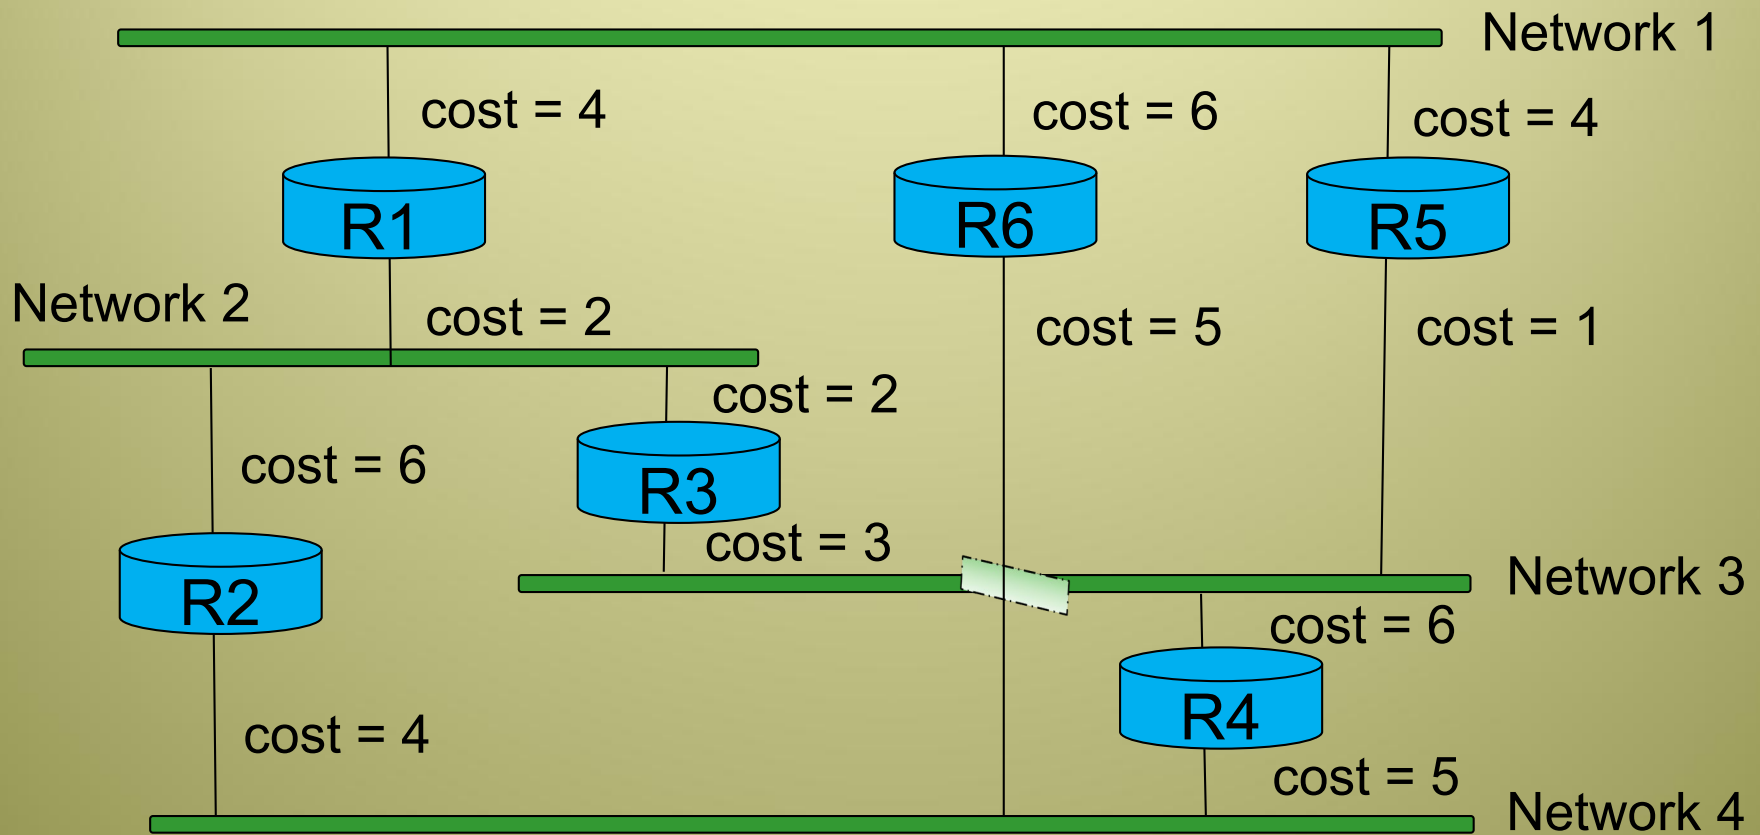
(a) stepwise construction
of spanning tree
(center: E)



(b) constructed spanning
tree

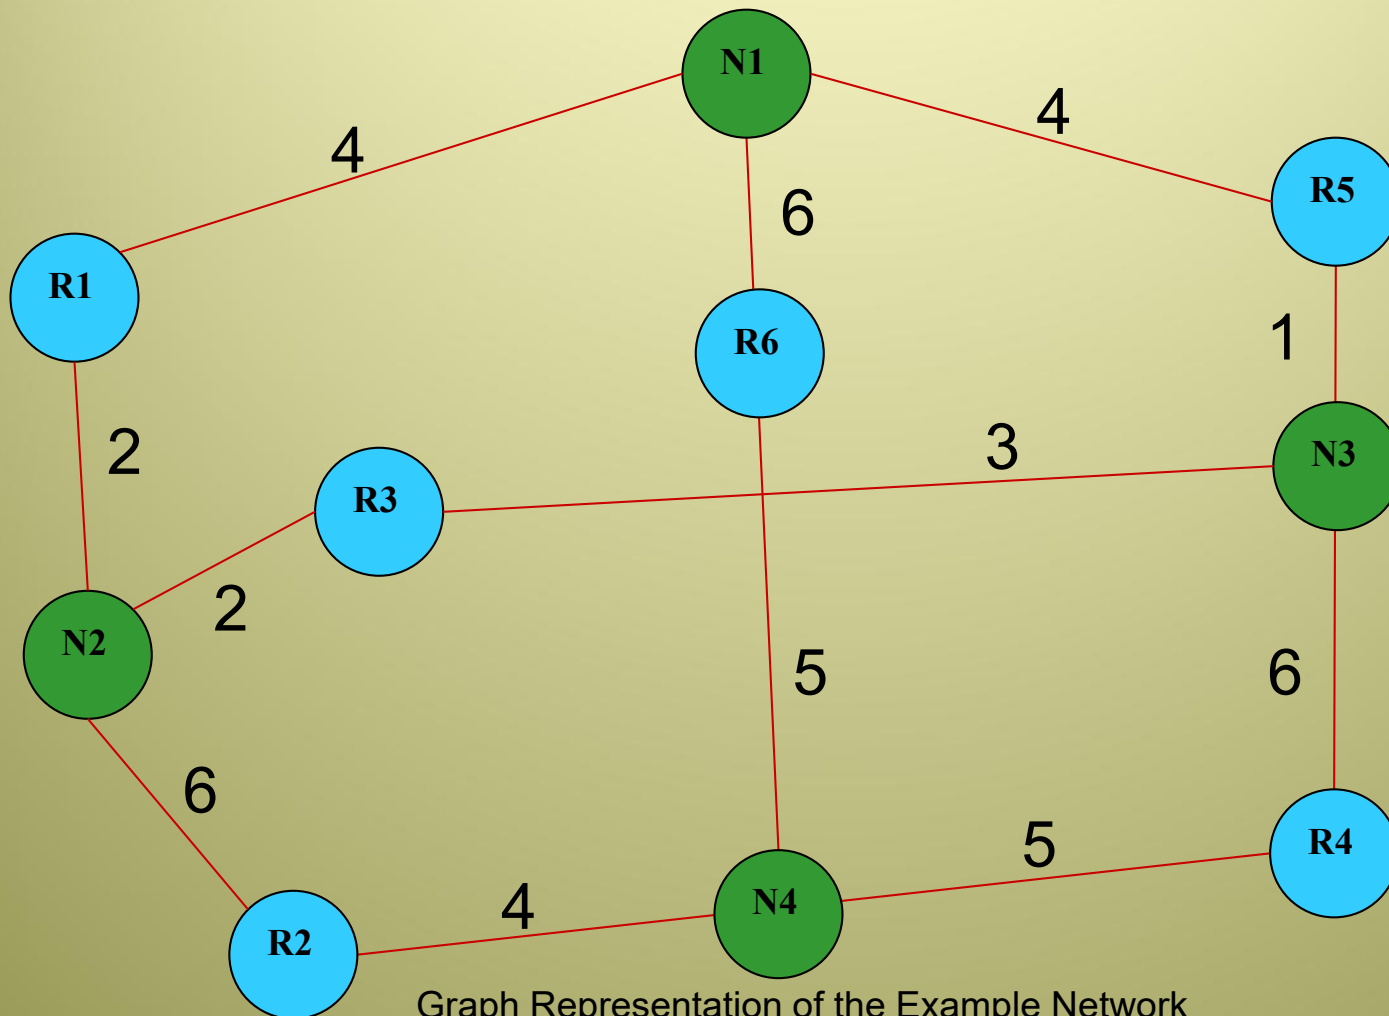
Spanning tree: creation

- ❖ If each link has an associated cost, the spanning tree can be constructed based on this cost



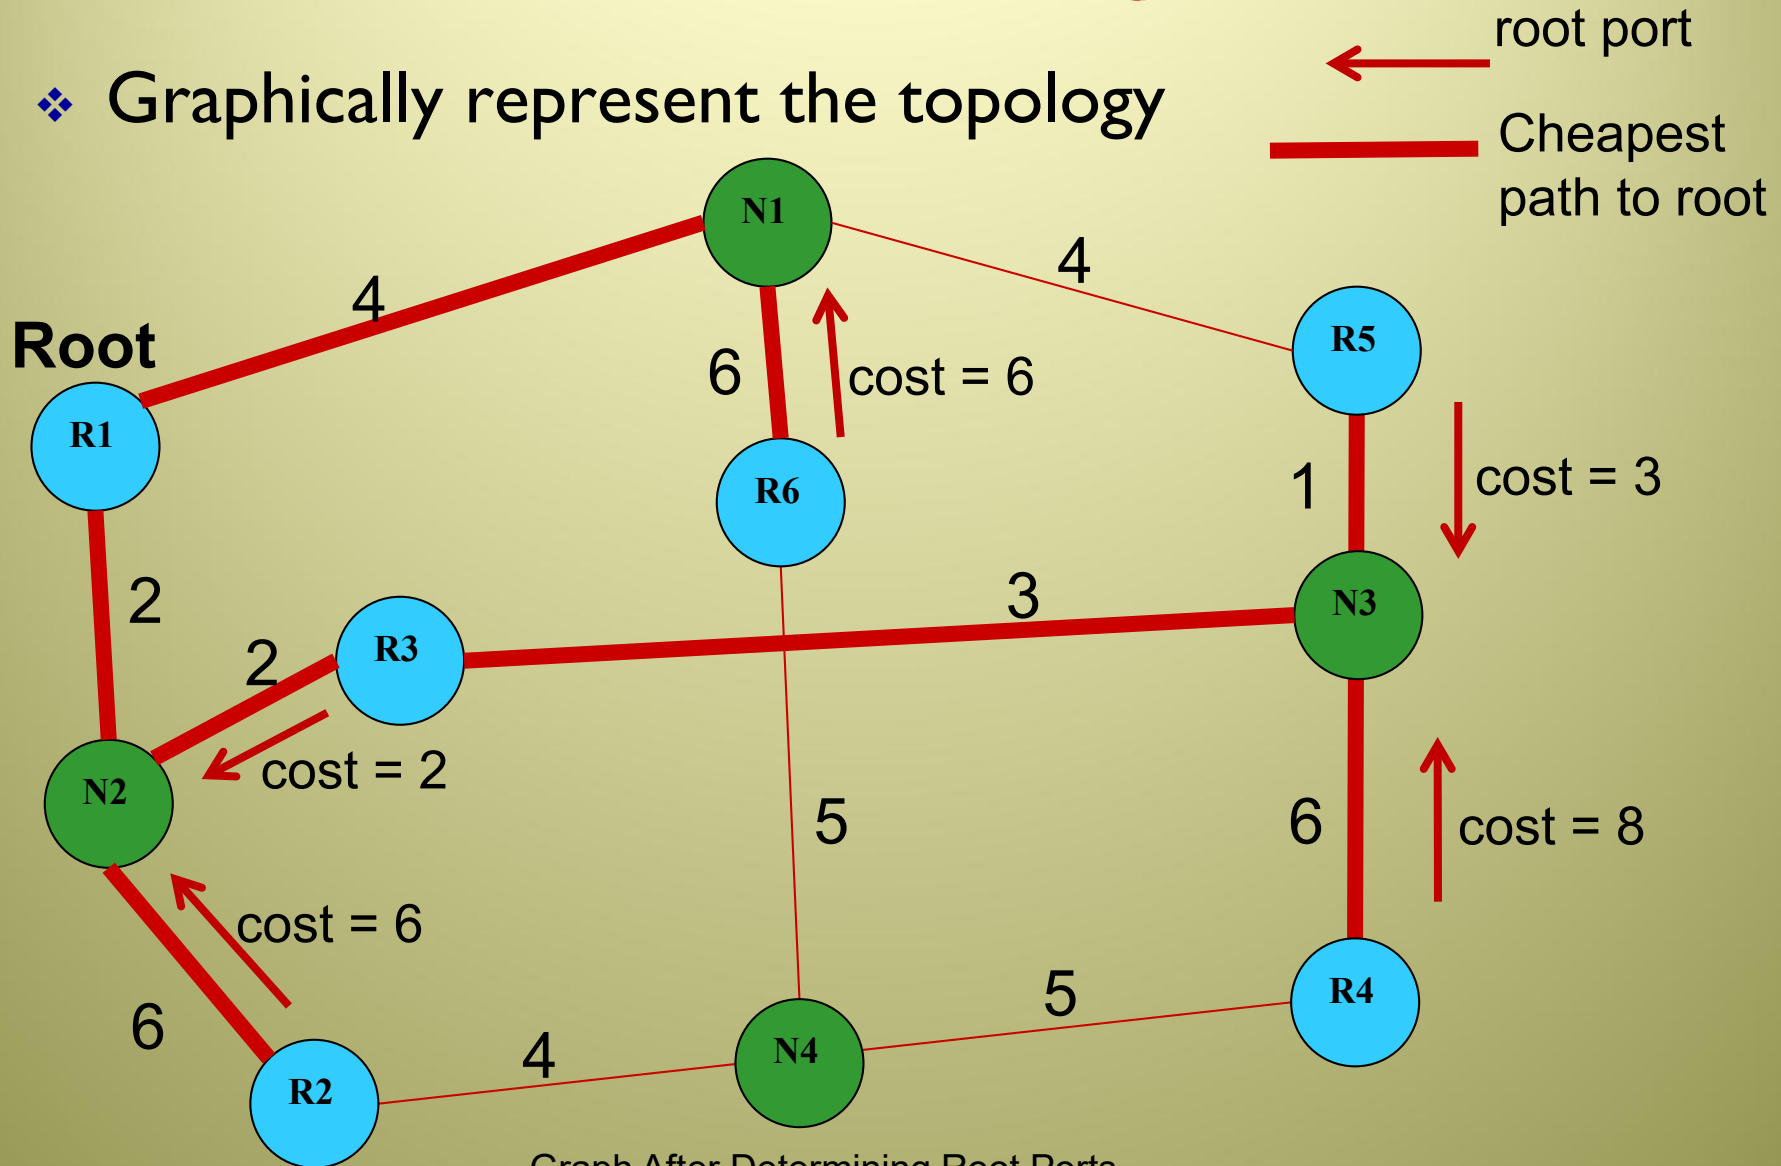
Spanning tree: creation

❖ Graphically represent the topology



Spanning tree: creation


- ❖ Graphically represent the topology



Graph After Determining Root Ports

Spanning tree: creation

❖ Construct the spanning tree



Spanning Tree of the Example Network

Spanning tree: creation

- ❖ network can hence be setup as follows (dashed links represent routers not active in flooding)

